

## Section 8

```
library(coda)
library(MASS)
library(tidyverse)
```

### Metropolis Algorithm

To generate sample  $s + 1$  of a Metropolis MCMC sampler given (possibly) unnormalized density  $p(\theta)$ :

1. Propose a new sample  $\theta_*$  given the old sample  $\theta_s$  from a symmetric distribution
2. If  $p(\theta_*) > p(\theta_s)$ , then set  $\theta_{s+1}$  equal to  $\theta_*$  and go to the next iteration
3. If  $p(\theta_*) < p(\theta_s)$ , then
  - a. Generate a random number  $r$  from  $\text{uniform}(0, 1)$
  - b. If  $r < \frac{p(\theta_*)}{p(\theta_s)}$ , set  $\theta_{s+1}$  equal to  $\theta_*$  and go to the next iteration
  - c. Otherwise set  $\theta_{s+1}$  equal to  $\theta_s$

### Little wheels

Assume you are given the following (incredibly inefficient) code and told to use it to sample from a normal distribution

```
# pdf we want to sample
p = function(theta) {
  dnorm(theta, 1.0, 2.0)
}

metropolis = function(theta_s) {
  # Function should return the next state
  # in the Markov chain given the current state, theta_s!

  theta_p = rnorm(1, theta_s, 1.0)

  if(p(theta_p) > p(theta_s)) {
    return(theta_p)
  } else {
    r = runif(1)
    if(r < p(theta_p) / p(theta_s)) {
      return(theta_p)
    } else {
      return(theta_s)
    }
  }
}

N = 10000
samples = rep(0, N)
samples[1] = 10.0
for(i in 1:(length(samples) - 1)) {
```

```
Samples[i + 1] = metropolis(samples[i])
}
```

We can look at our traceplots and effective sample size estimates with the coda package:

```
plot(as.mcmc(samples))
effectiveSize(samples)
```

## Medium wheels

Using the un-logged densities is numerically unstable. As an example of what can happen, compare the outputs of:

```
print(1.0e-100 * 1e-100, format = "e", digits = 20)
```

```
## [1] 9.999999999999999821e-201
```

```
print(1.0e-200 * 1e-200, format = "e", digits = 20)
```

```
## [1] 0
```

It is really common to need to evaluate numbers this small in a probabilistic model. For instance, a term like  $0.36^{300}$  might come up when evaluating a binomial pmf that models a basketball player's yearly shooting percentage. If we extend that to maybe three years worth of shots-made,  $(0.36^{900})$ , we'll see that evaluates to zero. The trick to avoid this is working on the log scale.

We want our metropolis algorithm to work on a log scale too. Because log is a monotonic increasing function, we can just take the log of the conditions in steps 2 and 3 above and get our new algorithm:

2. If  $\log p(\theta_*) > \log p(\theta_s)$ , then set  $\theta_{s+1}$  equal to  $\theta_*$  and go to the next iteration
3. If  $\log p(\theta_*) < \log p(\theta_s)$ , then
  - a. Generate a random number  $r$  from  $\text{uniform}(0, 1)$
  - b. If  $\log(r) < \log p(\theta_*) - \log p(\theta_s)$ , set  $\theta_{s+1}$  equal to  $\theta_*$  and go to the next iteration
  - c. Otherwise set  $\theta_{s+1}$  equal to  $\theta_s$

Rewrite the code above to work on the log scale and convince yourself it is working.

```
logp = function(theta) {
  dnorm(theta, 1.0, 2.0, log = TRUE)
}
metropolis = function(theta_s) {
  # Use logp, *not* log(p(...))
  theta_p = rnorm(1, theta_s, 1.0)

  if(logp(theta_p) > logp(theta_s)) {
    return(theta_p)
  } else {
    r = runif(1)
    if(log(r) < logp(theta_p) - logp(theta_s)) {
      return(theta_p)
    } else {
      return(theta_s)
    }
  }
}
N = 10000
samples = rep(0, N)
```

```

samples[1] = 100.0
for(i in 1:(length(samples) - 1)) {
  samples[i + 1] = metropolis(samples[i])
}

```

## Big wheels

Because you are enterprising young statisticians, you want to sample a multidimensional distribution. You can use the function `mvrnorm` to sample from a multivariate proposal distribution like so:

```

library(MASS)
mvrnorm(1, c(1.0, 2.0), matrix(c(1.0, 0.5, 0.5, 2.0), nrow = 2))

```

```
## [1] 1.817168 2.727315
```

This is a sample from:

$$N\left(\begin{bmatrix} 1 \\ 2 \end{bmatrix}, \begin{bmatrix} 1 & 0.5 \\ 0.5 & 2 \end{bmatrix}\right)$$

Because you are enterprising young software engineers, you want to write a function that does the Metropolis sampling for you, handles the burnin, and makes it easy to try different proposal covariances.

Because this is PSTAT115, you get to do that now:

```

logp = function(theta) {
  sum(dnorm(theta, c(1.0, 2.0), c(2.0, 3.0), log = TRUE))
}

metropolis = function(theta_s, cov) {
  # Note: theta_s is a vector now!
  theta_p = mvrnorm(1, theta_s, cov)

  if(logp(theta_p) > logp(theta_s)) {
    return(theta_p)
  } else {
    r = runif(1)
    if(log(r) < logp(theta_p) - logp(theta_s)) {
      return(theta_p)
    } else {
      return(theta_s)
    }
  }
}

# theta_0 is your initial parameter guess
# burnin is number of burnin samples
# maxit are the number of samples you generate (post burnin)
# cov is the proposal covariance matrix
rw_metrop_multi = function(theta_0, burnin, maxit, cov) {
  samples = matrix(0, ncol = length(theta_0), nrow = burnin + maxit)

  samples[1,] = theta_0
  for(i in 1:(nrow(samples) - 1)) {
    samples[i + 1,] = metropolis(samples[i,], cov)
  }
  samples[burnin:(burnin + maxit),]
}

```

```
}
samples = rw_metrop_multi(c(10.0, 10.0), 1000, 10000, matrix(c(1.0, 0.0, 0.0, 2.0), nrow = 2))
```

Check that the sampler is producing results you trust, because we are about to work on a non-trivial problem!

## Model and Simulated data

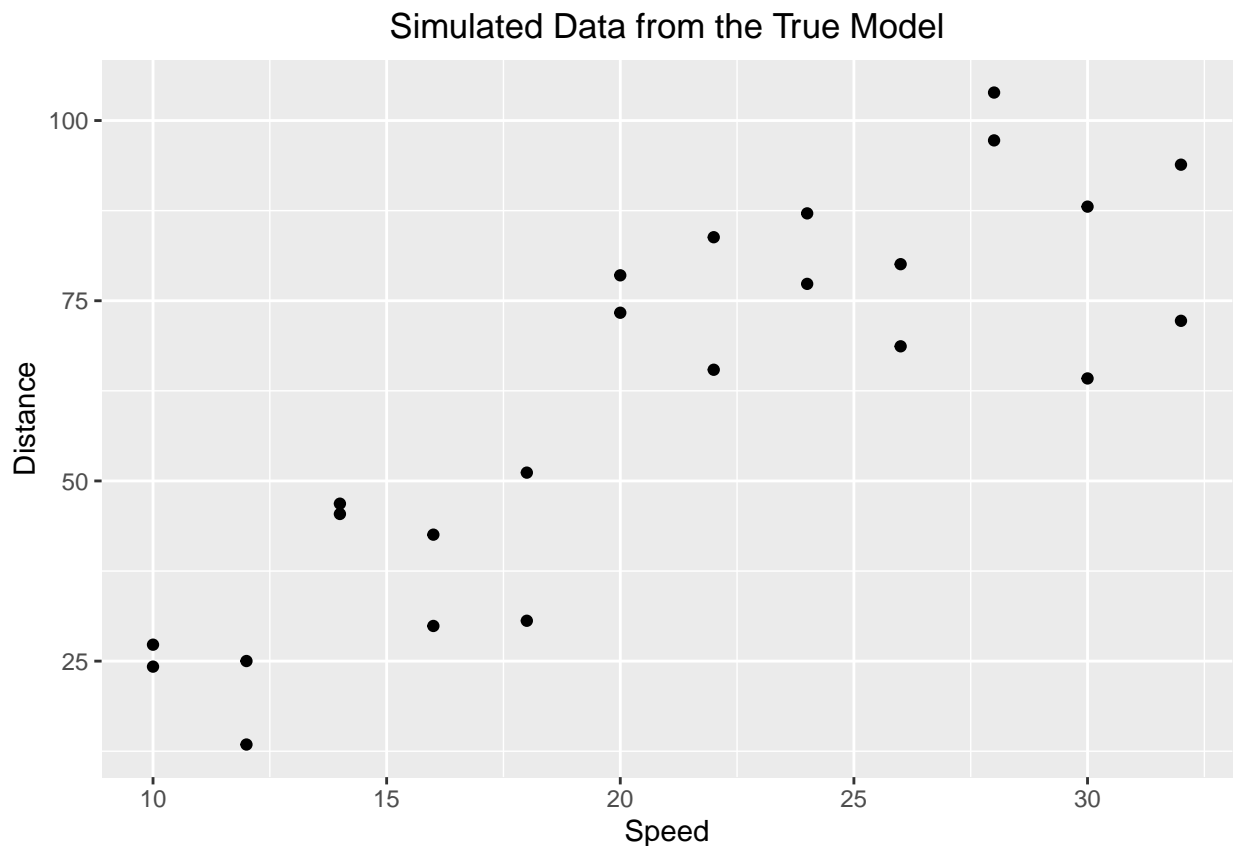
Let us consider a model of the speed of cars and the distances taken to stop. Assuming the true model is

$$dist = -8 + 3.5 * speed + \epsilon, \text{ where } \epsilon \sim N(0, 15^2)$$

And now we generate simulated data from the true model and try to estimate the parameters using the Bayesian approach.

```
true_beta0 <- -8
true_beta1 <- 3.5
x <- rep(seq(10, 32, by = 2), each = 2)
set.seed(10)
y <- true_beta0 + true_beta1 * x + rnorm(length(x), 0, 15)

data.frame(x = x, y = y) %>% ggplot() + geom_point(aes(x = x, y = y)) +
  xlab("Speed") + ylab("Distance") + ggtitle("Simulated Data from the True Model") +
  theme(plot.title = element_text(hjust = 0.5))
```



Now we try to propose a simple linear model and estimate the parameters using Metropolis-Hastings algorithm.

Our model is

$$dist = \beta_0 + \beta_1 * speed + \epsilon, \text{ where } \epsilon \sim N(0, 15^2).$$

Notice that here we assume we know the variance for the error term, so we can focus on the estimation of  $\beta_0$  and  $\beta_1$ .

## Function for log\_posterior

From the lecture materials we know it is often more stable working with log-scale. So here we write a function for the log\_posterior. You only need to consider the log for the essential parts in the posterior distribution.

```
logp <- function(beta){  
  return(sum(dnorm(y - beta[1] - beta[2] * x, mean = 0, sd = 15, log = TRUE)))  
}
```

## Call our Metropolis sampler

```
samples <- rw_metrop_multi(c(0, 0), 1000, 20000, cov = matrix(c(30.0, 0.0, 0.0, 1.0), nrow=2))  
plot(samples[,1], samples[,2])
```

## Covariance Matrix for the Proposal Distribution

We can actually play around with the covariance matrix in the proposal distribution. Negative correlation between  $\beta_0$  and  $\beta_1$  tend to provide better sampling results.