

Lecture 7: Markov Chain Monte Carlo

Professor Alexander Franks

2020-11-16

Announcements

Homework 4 out.

- 3 Q's

Monte Carlo estimation

- $\bar{\theta} = \sum_{s=1}^S \theta^{(s)} / S \rightarrow \mathbb{E}[\theta | y_1, \dots, y_n]$

- $\sum_{s=1}^S (\theta^{(s)} - \bar{\theta})^2 / (S - 1) \rightarrow \text{Var}[\theta | y_1, \dots, y_n]$

- $\# (\theta^{(s)} \leq c) / S \rightarrow \Pr(\theta \leq c | y_1, \dots, y_n)$

- the α -percentile of $\{\theta^{(1)}, \dots, \theta^{(S)}\} \rightarrow \theta_\alpha$


$$\int \theta P(\theta | y) d\theta$$

$$\int (\theta - \mathbb{E}[\theta | y])^2 P(\theta | y) d\theta$$

$$\int I[\theta \leq c] P(\theta | y) dy$$

Sampling from the posterior distributions

- The Monte Carlo methods we discussed previously assumed we could easily get samples from the posterior, e.g. with `rnorm`
- In general, sampling from a general probability distribution is hard
- Want to call `rcomplicatedistribution()` but don't have it
 - Inversion and rejection can work well for low dimensional posteriors
- In high dimensions, these approaches aren't sufficient
 - Near impossible to find good proposal distributions that envelope target
 - Or rejection rate is extremely high



Markov Chain Monte Carlo

- We want independent random samples, $\theta^{(s)}$ from $p(\theta \mid y_1, \dots, y_n)$
- But there is no good way to get independent samples
- Alternative, create a sequence of correlated samples with the correct **limiting** distribution
- Markov Chain Monte Carlo gives us a way to generate correlated samples from a distribution

time series

$\theta^{(s)} \sim \text{MCMC (correlated samples)}$

As long as
 $E[\theta^{(s)}] = E[\theta \mid y]$

$$\left[E\left[\frac{1}{n} \sum \theta^{(s)}\right] \xrightarrow{\text{lim}} E[\theta \mid y] \right] \text{ (linearity of expectation)}$$

*Correlation doesn't hurt us.
 in this particular regard.*

Monte Carlo Error

- Reminder: $\bar{\theta} = \sum_{s=1}^S \theta^{(s)} / S$ and S is the number of samples.

- If the samples are independent,

Monte Carlo Error

$$\text{Var}(\bar{\theta}) = \frac{1}{S^2} \sum_{s=1}^S \text{Var}(\theta^{(s)}) = \frac{\text{Var}(\theta \mid y_1, \dots, y_n)}{S} = O\left(\frac{1}{S}\right)$$

- If the samples are positively correlated,

$$\text{Var}(\bar{\theta}) = \frac{1}{S^2} \sum_{s,t} \text{Cov}(\theta^{(s)}, \theta^{(t)}) > \frac{\text{Var}(\theta \mid y_1, \dots, y_n)}{S}$$

$$\frac{1}{S^2} \left(\sum_{i=1}^S \text{Var}(\theta^i) \right) + \sum_{i \neq j} \text{Cov}(\theta^i, \theta^j)$$

0 when indep.

- MCMC methods have higher Monte Carlo error due to positive dependence between samples.
- Hope to minimize dependence and hence MC error

Basics of Markov Chains

Markov Chains: Big Picture

- For standard Monte Carlo, we make use of the law of large number to approximate posterior quantities
- The law of large numbers can still apply to random variables that are not independent
- We have a sequence of random variables indexed in time, θ_t
- We'll be using a discrete-time Markov Chain: $t \in 0, 1, \dots, T$
- The observations, $\theta^{(t)}$ can be discrete or continuous ("discrete-state" or "continuous-state" Markov Chain)

Discrete-state Markov Chains

- Let $\theta^{(t)} \in 1, 2, \dots, M$ be the state space for the Markov Chain
- A sequence is called a markov chain if

$$Pr(\theta^{(t+1)} \mid \theta^{(t)}, \theta^{(t-1)} \dots \theta^{(1)}) = Pr(\theta^{(t+1)} \mid \theta^{(t)})$$


for all $t \geq 0$

- The **Markov property**: given the entire past history, $\theta^{(1)}, \dots, \theta^{(t)}$, the most recent $\theta^{(t+1)}$ depends only on the immediate past, $\theta^{(t)}$

Memory of 1 time period

The Transition Matrix

- Define $q_{ij} = \Pr(\theta^{(t+1)} \mid \theta^{(t)})$ is the transition probability from state i to state j
- The $M \times M$ matrix $Q = (q_{ij})$ is called the *transition matrix* of the Markov Chain



3-state example

$$\begin{array}{c} \text{From} \\ \underline{Q} = \end{array} \begin{array}{c} \text{To} \\ \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix} \end{array}$$

The Transition Matrix

3-state example

$$Q = \begin{bmatrix} q_{11} & q_{12} & q_{13} \\ q_{21} & q_{22} & q_{23} \\ q_{31} & q_{32} & q_{33} \end{bmatrix}$$

$$q_{11} + q_{12} + q_{13} = 1$$

- The rows of the transition matrix sum to 1
- Note: $Q^n = (q_{ij}^{(n)})$ is the probability of transitioning from i to j in n steps

$$(Q \times Q)_{ij} = \text{Prob. of going to } j \text{ in 2 steps, given start in } i.$$

- A Markov Chain is **regular** if Q^n has strictly positive entries for some value of $n > 1$
- A Markov Chain is **irreducible** if for any two states i and j it is possible to go from i to j in a finite number of steps (with positive probability)

"Graph is connected"



Reducible

The limiting distribution

- A regular, irreducible Markov chain has a **limiting probability distribution**
- Describes the long-run fraction of time does the Markov Chain spend in each state (in the long run)
 - *Does not* depend on where the chain starts
- Let $\pi = (\pi_1, \dots, \pi_M)$ be a row vector of probabilities associated with each state, such that $\sum_{i=1}^M \pi_i = 1$
 - The limiting distribution converges to π , which is said to be **stationary** because ~~$\pi Q = \pi$~~ $Q^T \pi = \pi$ ^{long-run island probabilities}
 - If you sample from the limiting distribution and then transition, the result is still distributed according to the limiting distribution

e.g. $\pi = (1/4, 1/4, 1/2)$

Markov Chain Example

- Sociologists often study social mobility using a Markov chain.
- In this example, the state space is {low income, middle income, and high income} of families
- Let \mathbf{Q} be the transition matrix from parents income to childrens income

Parents $\mathbf{Q} =$

	Lower	Middle	Upper
Lower	0.40	0.50	0.10
Middle	0.05	0.70	0.25
Upper	0.05	0.50	0.45

Child

Multi-step Transition Probabilities

2-step transition probabilities

$$Q^2 = Q \times Q = \begin{matrix} & \begin{matrix} L & M & V \end{matrix} \\ \begin{matrix} \text{parents} \end{matrix} & \begin{vmatrix} 0.1900 & 0.6000 & 0.2100 \\ 0.0675 & 0.6400 & 0.2925 \\ 0.0675 & 0.6000 & 0.3325 \end{vmatrix} \end{matrix}$$

4-step transition probabilities

$$Q^4 = Q^2 \times Q^2 = \begin{matrix} \text{4th Generation} & \begin{vmatrix} 0.0908 & 0.6240 & 0.2852 \\ 0.0758 & 0.6256 & 0.2986 \\ 0.0758 & 0.6240 & 0.3002 \end{vmatrix} \end{matrix}$$

Multi-step Transition Probabilities

4-step transition probabilities

$$\mathbf{Q}^4 = \mathbf{Q}^2 \times \mathbf{Q}^2 = \begin{vmatrix} 0.0908 & 0.6240 & 0.2852 \\ 0.0758 & 0.6256 & 0.2986 \\ 0.0758 & 0.6240 & 0.3002 \end{vmatrix}$$

8-step transition probabilities

$$\mathbf{Q}^8 = \mathbf{Q}^4 \times \mathbf{Q}^4 = \begin{vmatrix} 0.0772 & 0.6250 & 0.2978 \\ 0.0769 & 0.6250 & 0.2981 \\ 0.0769 & 0.6250 & 0.2981 \end{vmatrix}$$

Converging to limiting distribution.

Idea of MCMC

The limiting distribution

Design Q
so $\pi = P(\theta|y)$

$$Q^\infty = \mathbf{1}\pi =$$

π_1	π_2	π_3
π_1	π_2	π_3
π_1	π_2	π_3

π is limiting distribution,

```
Q <- matrix(c(0.4, 0.05, 0.05,  
              0.5, 0.7, 0.5,  
              0.1, 0.25, 0.45),  
            ncol=3)
```

```
p <- eigen(t(Q))$vectors[, 1]  
stationary_probs <- p/sum(p)  
stationary_probs
```

```
## [1] 0.07692308 0.62500000 0.29807692
```

```
stationary_probs %*% Q
```

```
##           [,1] [,2] [,3]  
## [1,] 0.07692308 0.625 0.2980769
```

$$Q^T \pi = \pi \quad (\text{defn of stationary})$$

limiting distn.

still in stationary distn.


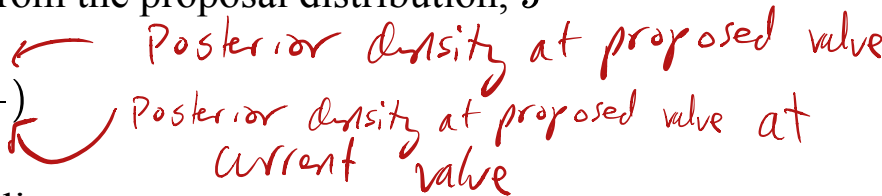
Markov Chain Monte Carlo

- Incredible idea: create a Markov Chain with the desired limiting distribution
- Want the limiting distribution to be the posterior distribution
- Unlike the previous examples, we will mostly work with *infinite* state space
- Instead of a transition matrix we have a transition kernel which is a conditional probability, $p(\theta^{(t+1)} \mid \theta^{(t)})$
- *Design* $p(\theta^{(t+1)} \mid \theta^{(t)})$ to have limiting distribution $p(\theta \mid y)$
 - If we run the random walk for long enough, $\theta^{(t)}$ will be distributed approximately according to $p(\theta \mid y)$
- The Metropolis algorithm tells us how to construct such a transition ~~matrix~~ *kernel*

Generalizing the rejection sampler

- Make a small tweak to the rejection sampler
- Sample from a proposal, $q(\theta)$, doesn't have to envelope $p(\theta | y)$!
- If $p(\theta | y) > 0$ then we need $q(\theta) > 0$ (same support)
- Unlike the rejection sampler, we never "throw out" samples
- Instead, at each iteration we have a choice:
 - ◦ Accept the new proposed sample
 - Or
 - ◦ Or **repeat** the previous sample again

Generalizing the rejection sampler

1. Initialize θ_0 to be the starting point for you Markov Chain
2. Choose a *proposal distribution*, $J(\theta^*)$  "Jump" density
 - Propose a candidate value for the next sample
 - Best performance if density is very similar to target
3. Generate the candidate θ^* from the proposal distribution, J
4. Compute $r = \min(1, \frac{p(\theta^*|y)}{p(\theta_t|y)})$ 
 - Generate a uniform random number $u \sim Unif(0, 1)$
 - If $u \leq r$ we accept θ^* as our next sample
 - Else $\theta_{t+1} \leftarrow \theta_t$ (we do not update the sample this time)

Intuition

- If $p(\theta^* | y) > p(\theta_t | y)$ accept with probability 1
 - The proposed sample has higher posterior density than the previous sample
 - Always accept if we increase the posterior probability density
- If $p(\theta^* | y) < p(\theta_t | y)$ accept with probability $r < 1$
 - Accept with probability less than 1 if probability density would decrease
 - Relative frequency of θ^* vs θ_t in our samples should be $\frac{p(\theta^*|y)}{p(\theta_t|y)}$

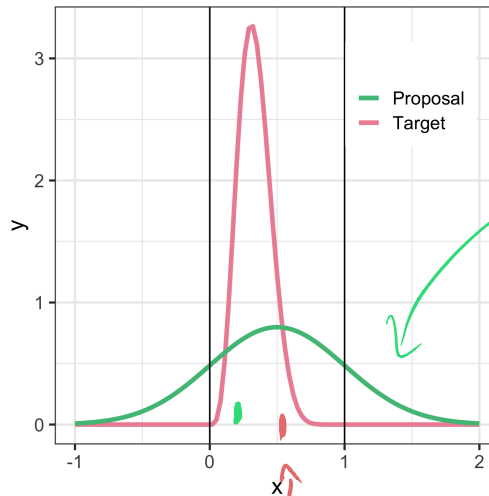
Independence Sampler

- The previous algorithm is known as an "Independence Sampler"
- Let $P(\theta | y)$ be a Beta(5, 10) posterior distribution
- Propose from a distribution $J(\theta^*) \sim N(0.5, 1)$

confusing
Name!

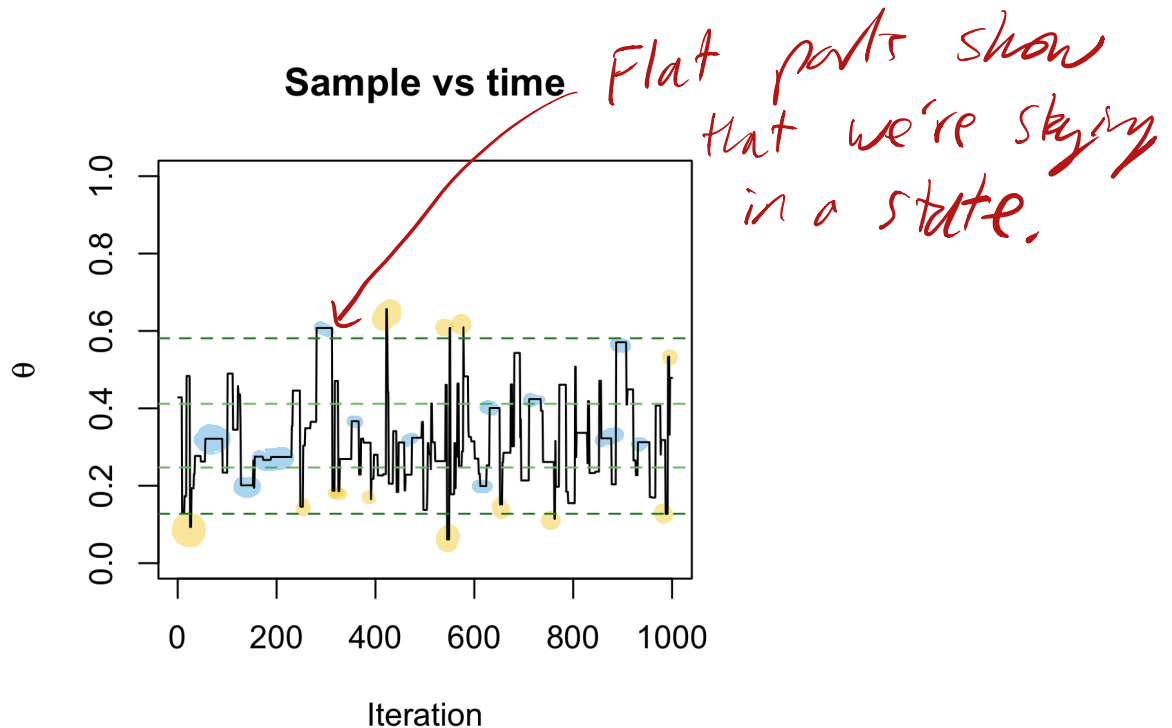
Samplers are not
independent.

$$\frac{P(\theta^* | y)}{P(\theta_0 | y)}$$




Proposal
(Doesn't envelope
target)

Independence Sampler



Note and source of confusion: samples are correlated over time for the "independence sampler".


Weighting by waiting



Where did the sampler get stuck? Where does it quickly leave?

Independence Sampler

Monte Carlo vs True



N = 100000 Bandwidth = 0.01059

The Metropolis Algorithm

- The Metropolis Algorithm generalizes the independence sampler
- Allow the proposal distribution to depend on the most recent sample

- Independence: $J(\theta^*)$, e.g. $\theta^* \sim N(0.5, 1)$

Does not depend on θ_t

- Metropolis: $J(\theta^* | \theta_t)$, e.g. $\theta^* \sim N(\theta_t, 1)$

- Independence sampler: "Independence" refers to the proposal being fixed (**not** independence samples)!
- Metropolis sampler: a "moving" proposal distribution

Proposal density changes in time.

The Metropolis Algorithm

1. Initialize θ_0 to be the starting point for you Markov Chain
2. Choose a proposal distribution, $J(\theta^* | \theta_t)$
 - Propose a candidate value for the next sample
 - Must have symmetry: $J(\theta^* | \theta_t) = J(\theta_t | \theta^*)$
3. Generate the candidate θ^* from the proposal distribution, J
4. Compute $r = \min(1, \frac{p(\theta^*|y)}{p(\theta_t|y)})$
5. Set $\theta_{t+1} \leftarrow \theta^*$ with probability r
 - Generate a uniform random number $u \sim Unif(0, 1)$
 - If $u < r$ we accept θ^* as our next sample
 - Else $\theta_{t+1} \leftarrow \theta_t$ (we do not update the sample this time)

Metropolis Algorithm

- Let $P(\theta \mid y)$ be a $\text{Beta}(5, 10)$ posterior distribution
- 1-d sampling: lets try sampling from the Beta using the Metropolis algorithm
- Initialize θ_0 to 0.9
 - Note that the probability of drawing a value larger than 0.9 from a $\text{Beta}(5, 10)$ is smaller than $1\text{e-}8$
 - Our initial value is far from the high posterior density
 - In the long run this won't matter
- Define transition kernel $J(\theta_{t+1} \mid \theta_t)$ as $\theta^* \sim N(\theta_t, \tau^2)$
 - How does choice of τ^2 effect performance of MC sampler?